





localisation: β_a vs its N2 null



KRAS twin: does a flagged cliff replicate?



—●— (a) observed $\beta_a \pm 1.96$ HC3 SE ■ (a) N2 null 2.5th-97.5th | (a) N2 99.5th pct —●— (b) R with 95% CI ■ (b) permutation chance level

KRAS-RAF1 vs KRAS-RAF1-RBD, $n=10868$ shared site pairs, $r=0.812$, OLS slope 0.644, sign agreement 0.952 against a 0.596 chance level, permutation $p=1e-04$. Panel (b): spec Sec.6 asks for a hexbin of ϵ_a vs ϵ_b over the shared site pairs. T08 tabulates that join's SUMMARY (n_{shared} , pearson_raw, ols_slope, sign agreement, R and its permutation chance level) but not the per-pair ϵ ; T10 carries per-assay ϵ , and joining two assays by hand is exactly the naive cross-assay join gate G1b bans, so the panel shows what T08 does tabulate: R at every threshold variant against its own permutation chance level; the measured summary statistics open this footer.

Panel (c): T11 tabulates the RATIO F_{spec} (and its noise-corrected form, with $\sigma^2 = 0.01545$ subtracted), not the three variance components separately, so the panel shows the ratio and not the stacked shared / specific / noise decomposition spec Sec.6 describes. "at bound" = $F_{\text{spec_at_boundary}}$, i.e. the estimate sits at 1.00 and the family carries no usable shared component.

Panel (d): T10's ICC_sitepair is ONE value per assay repeated on every site-pair row (verified: at most one distinct value per assay), so Sec.6's "ICC across amino-acid combinations" cannot be drawn from it; the per-assay value is shown instead and the count of site pairs and the largest number of amino-acid combinations behind it are annotated.